

# Introduction to R for data analysis

Peter Carbonetto

Research Computing Center and the Dept. of Human Genetics  
University of Chicago



# Aims of workshop

1. Work through a basic data analysis in R.
2. Understand how to import data from a CSV file into an R data frame.
3. Use standard tools to summarize & manipulate data frames.
4. Learn how to install & use R packages.
5. Use ggplot2 to create plots from data frames.
6. Learn through “live coding”—this includes learning from our mistakes!

# Our goal: Analyze Divvy data from 2016 & 2017

- Investigate bike sharing trends in Chicago.
- We will use data made available from Divvy:
  - ▷ [www.divvybikes.com/system-data](http://www.divvybikes.com/system-data)
- Much of the effort will be spent importing the data, inspecting the data, and preparing the data for analysis.
- Once we have carefully prepared the data, creating visualizations is (relatively) easy.

# It's your choice

Your may choose to . . .

- Use R on your laptop.
- Use RStudio on your laptop.
- Use R or RStudio on the RCC cluster.
- Pair up with your neighbour.
- Follow what I do on the projector.

**Note:** If you use the RCC cluster I'm assuming you know how to set up an interactive computing session with appropriate amount of compute time and memory, load R or RStudio, and display graphics (e.g., using ThinLinc).

# Software we will use today

1. **R**
2. R packages **readr**, **ggplot2** & **cowplot**.
3. **RStudio** (optional).

**Note:** I'm assuming you have already installed R and/or RStudio on your laptop, or you are using the RCC cluster.

# Outline of workshop

1. Initial setup.
2. Load & prepare the Divvy station data.
3. Load & prepare the Divvy trip data.
4. Create a map of the Divvy stations.
5. Create plots comparing bike sharing in 2016 & 2017.

# Initial setup

- WiFi
- Power outlets
- Pace, questions (e.g., keyboard shortcuts).
- Help.

# Download or “clone” git repository

Download the workshop packet to your computer.

- Go to: **<https://github.com/rcc-uchicago/R-intro-divvy>**
- To download, click the green “**Clone or download**” button.

Or, if you have **git**, run this command:

```
git clone https://github.com/rcc-uchicago/  
R-intro-divvy.git
```

(Note the URL in the git command should not contain any spaces.)

- **Note:** If you are using the RCC cluster, also download workshop packet on the cluster.



# What's included in the workshop packet

1. **analysis**: directory where the code and data for our analyses will be stored.
2. **slides.pdf**: the workshop slides.
  - ▷ This PDF is useful for copying & pasting code from the slides. (Although I have found that you should only copy one line at a time.)
  - ▷ The PDF is included in the workshop packet.
  - ▷ You can also view the PDF by clicking the “slides.pdf” item on the GitHub webpage.

# Set up your R environment

- Launch R or RStudio.
- *We will run all the code from the “**analysis**” folder.*
- To change your working directory:
  - ▷ In R, use `setwd()` function.
  - ▷ In RStudio, select **Session > Set Working Directory > Choose Directory...**

Before continuing, check that you have the right working directory:

```
getwd()    # Should end with "analysis".
```

# Run `sessionInfo()`

Check the version of R that you are using:

`sessionInfo()`

If you are using an older version of R (version 3.3 or less), I strongly recommend upgrading to the latest version. *Some of the examples may not work in older versions of R.*

# Check your R environment

The R environment is where all variables and functions are stored and accessed. You should start with an empty environment:

```
ls()
```

If you see names of objects listed, it means your environment is not empty, and you should restart R with a clean environment.

- Do `rm(list = ls())`.
- Or, in RStudio, go to **Session > Restart R**.

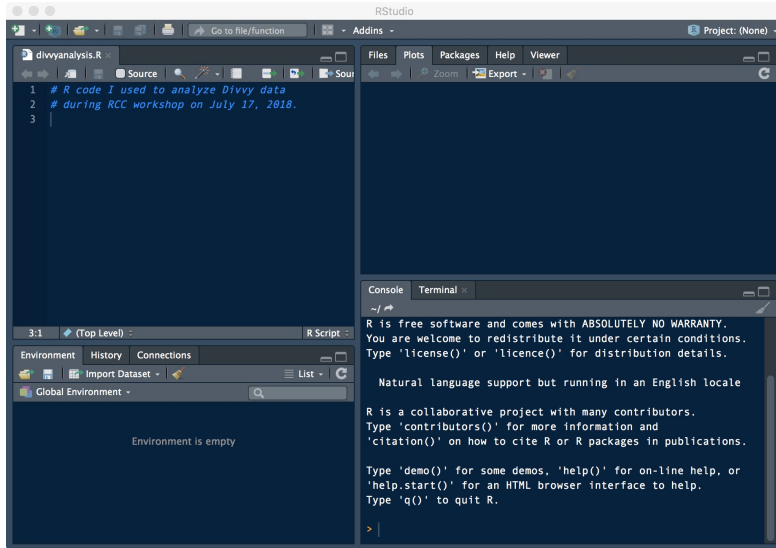
# Creating a file to keep track of your analysis code

- In RStudio, select **File > R Script**.
- Alternatively, use your favourite editor.
- Add some comments to the title to remind yourself what this file is for, e.g.,

```
# Some of the R code I wrote during the RCC  
# workshop on August 16, 2018.
```

- Save the file in the “**analysis**” folder. Name the file whatever you’d like (e.g., `divvyanalysis.R`).

# The Console is the “brains” of RStudio



# Download the Divvy data

- Disk space required: at least **2 GB**.
- Download the 2016 & 2017 data files from here:
  - ▷ **[www.divvybikes.com/system-data](http://www.divvybikes.com/system-data)**
- Download them to the “**analysis**” folder.
- You should have 4 ZIP files:

Divvy\_Trips\_2016\_Q1Q2.zip

Divvy\_Trips\_2016\_Q3Q4.zip

Divvy\_Trips\_2017\_Q1Q2.zip

Divvy\_Trips\_2017\_Q3Q4.zip

- Decompress (“unzip”) all of these files.

# Check that you have all the files

After unzipping, you should have **15** CSV files.

```
sys.glob ("*.CSV")
```

*If you don't see all 15 CSV files, you have not successfully downloaded and/or unzipped all the files in the "data" directory.*



# Outline of workshop

1. Initial setup.
2. **Load & prepare the Divvy station data.**
3. Load & prepare the Divvy trip data.
4. Create a map of the Divvy stations.
5. Create plots comparing bike sharing in 2016 & 2017.

# Eyeball the station data in the CSV file

- Open the CSV file **Divvy\_Stations\_2017\_Q1Q2.csv** in RStudio, or in your favourite text editor (e.g., Notepad in Windows, TextEdit on Mac).
- CSV is a simple and commonly used data format.
- It is easily read into R, and read by humans.
- Each line stores an item (station).
- The first line is a special line called the “header”.
- Entries in each line (table row) are separated by commas.

# Import the station data into R

Load the most up-to-date station data into an R “data frame”:

```
stations<-read.csv("Divvy_Stations_2017_Q3Q4.csv",  
                   stringsAsFactors = FALSE)
```

This will define a new object, “stations”, in your environment:

```
ls()
```

It is a “data frame” object:

```
class(stations)
```

*What does “read.csv” do, and what is a “data frame”?* R has detailed documentation:

```
help(read.csv)
```

```
help(data.frame)
```

# Inspect the station data

Run these commands to start inspecting the station data:

```
nrow(stations)
ncol(stations)
head(stations)
tail(stations)
summary(stations)
```

Inspect the data in more detail:

```
sapply(stations, class)
object.size(stations)
print(object.size(stations), units = "Kb")
```

*What do we learn about the station data from running these commands? Does this reveal any issues with the data?*

## Take a closer look at the “dpcapacity” column

Create a new object containing only the “dpcapacity” column:

```
x <- stations$dpcapacity
```

Run these commands to take a closer look at the “dpcapacity” column:

```
class(x)
```

```
length(x)
```

```
summary(x)
```

```
table(x)
```

*Did we gain any additional insight from running these commands?*

# Selecting rows & columns

Select first 4 rows of "name" column:

```
stations$name[1:4]  
stations[1:4, 2]  
stations[1:4, "name"]
```

Select first 4 rows and multiple columns:

```
stations[1:4, c(2, 3, 6)]  
stations[1:4, c("name", "city", "dpcapacity")]
```

Getting the row and column names:

```
colnames(stations)  
rownames(stations)
```

## Take an even closer look at “dpcapacity”

It is interesting that a few of the Divvy bike stations are much larger than the others, whereas others have no docks. Where are these stations?

```
subset(stations, dpcapacity == 0)
subset(stations, dpcapacity >= 40)
```

Alternatively, we can sort the table rows, then inspect the top and bottom rows:

```
rows <- order(stations$dpcapacity, decreasing=TRUE)
stations2 <- stations[rows,]
head(stations2)
tail(stations2)
```

*How were the rows originally ordered in stations?*

## Take a closer look at the “city” column

Above we inspected *numeric* data. Next's, let's look at an example of non-numeric data.

```
x <- stations$city  
class(x)  
summary(x)
```

The summary is not very useful here! The key is to convert to a “factor” (categorical variable):

```
x <- factor(stations$city)  
class(x)  
summary(x)
```

*Did you discover an issue with the data from running these commands?*



# Improving the “city” column

Let's fix the problem we found earlier. First, select the offending rows of the table:

```
rows <- which(stations$city == "Chicago ")
```

Fix the “city” column by overwriting the “city” entries in the selected rows:

```
stations[rows, "city"] <- "Chicago"  
summary(stations$city)
```

The “city” column is more useful if it is a factor, so let's convert the column directly inside the data frame:

```
stations$city <- factor(stations$city)  
summary(stations$city)
```

# What is a “factor”?

Factors are often very useful in data analyses. Let's take a deeper look at what a factor *is*.

```
x <- stations$city  
attributes(x)  
unclass(x)
```

From the `unclass(x)` call, we see that a factor is really just an integer with values 1, 2, 3, *etc.*, with which each integer value is associated with a *label* (e.g., “Chicago”, “Evanston”).

# Save your code & session state

It is important to periodically save:

1. your code,
2. the state of your R environment.

To save your environment, go to **Session > Save Workspace As...** in RStudio, or run this code:

```
save.image("divvyanalysis.RData")
```

Later, to restore your environment in a new session, select **Session > Load Workspace...** in RStudio, or run this code:

```
load("divvyanalysis.RData")
```

# Main concepts covered so far

- The R environment & working directory.
- Read a data frame from a text (CSV) file.
- Tools to inspect a data frame.
- Manipulate a data frame.
- Factors = categorical variables.
- Selecting rows & columns.
- Order rows of a data frame.
- Save state of R environment.

# Outline of workshop

1. Initial setup.
2. Load and prepare the Divvy station data.
3. **Load and prepare the Divvy trip data.**
4. Create a map of the Divvy stations.
5. Create a scatterplot comparing bike sharing activity in 2016 and 2017.

# Import the Divvy trip data into R

Previously, we used `read.csv` to import station data into R. Let's now use `read.csv` to load the trip data from the 4th quarter of 2017:

```
trips <-  
  read.csv("Divvy_Trips_2017_Q4.csv",  
           stringsAsFactors = FALSE)
```

You may find that this command took longer to run than before. Consider that the trips data is much larger:

```
nrow(trips)  
ncol(trips)  
print(object.size(trips), units = "Mb")
```

This gives an opportunity to demonstrate a faster method implemented in a *package*.

# Import Divvy trip data using readr (optional)

Install the **readr** package from CRAN:

```
install.packages("readr")
```

Load the functions from the package into your R environment:

```
library(readr)
```

Let's use the `read_csv` function from this package:

```
trips <- read_csv("Divvy_Trips_2017_Q4.csv")
```

**Note:** `read_csv` is similar to `read.csv`, but not the same.

- *How much faster is `read_csv`?*

# Import Divvy trip data using `readr` (optional)

The `read_csv` output is *not* a data frame—it is a “tibble”.

```
class(trips)
```

Typically, I convert it to a data frame:

```
class(trips) <- "data.frame"
```

- For more on tibbles, see:
  - ▷ <http://r4ds.had.co.nz>
- The **readr** package has many other features not covered here.
- Another fast method is `fread` from the **data.table** package.



# More on packages in R

“Vignettes” are a great way to learn about a package:

```
vignette(package = "readr")  
vignette("readr")
```

- CRAN is the official package source:

- ▷ <https://cran.r-project.org>.

- Other good places to find packages:

- ▷ Bioconductor
  - ▷ GitHub.

- *What packages are already installed?*

```
rownames(installed.packages())
```

- *Where do the packages live?* `.libPaths()`

- *How to learn more about a package?*

```
help(package=readr)
```

# A first glance at the trips data

Let's use some of the same commands we used earlier to quickly get an overview of the trip data:

```
nrow(trips)
ncol(trips)
head(trips)
summary(trips)
```

Unfortunately, the summary command isn't particularly informative for many of the columns.

- *What columns should we convert to factors?*

# Convert “gender” to a factor

Let's start by converting the “gender” column to a factor:

```
trips$gender <- factor(trips$gender)
summary(trips$gender)
levels(trips$gender)
```

We observe that many gender entries are missing.

# “Missing” data

- In R, “missing data” should always be assigned the special value `NA` (“not available” or “not assigned”).
- Many functions in R will correctly handle missing data as long as they are encoded as `NA`.
- The `read_csv` function from the `readr` package is “smart” enough to figure out that blank entries in the CSV file should be converted to `NA`.

# Convert “station” columns to factors

It is also useful to convert the “from station” column to a factor:

```
summary(trips$from_station_name)
trips$from_station_name <-
  factor(trips$from_station_name)
summary(trips)
```

The summary is now more informative.

## A note about dates & times

- `summary(trips)` is also not useful for the dates & times.
- Processing dates & times is more complicated.
- See `help(strptime)` and the **lubridate** package.

# Preparing data is tedious

Data preparation is sometimes >90% of the effort!

- *Many analysis mistakes are due to poor data preparation.*

Common issues include:

- Formatting mistakes in CSV file.
- Converting table columns to the appropriate data type.
- Entry inconsistencies (e.g., additional spaces).
- Missing data.
- Many other examples of Poor Practices in recording data.

(And we haven't yet dealt with merging data from multiple files—this usually creates more headaches!)

# Moving beyond data preparation

- So far, we have illustrated a few of the challenges of working with large tabular data sets (“data frames”).
- In order to proceed to fun stuff, I’ve automated the data preparation steps by writing a *function* in R to import and merge all the Divvy data into a single data frame.



# Outline of workshop

1. Initial setup.
2. Load & prepare the Divvy station data.
3. Load & prepare the Divvy trip data.
4. **Create a map of the Divvy stations.**
5. Create plots comparing bike sharing in 2016 & 2017.

# Clean up your environment

Here, we will begin a new analysis, so let's refresh our environment:

```
rm(list = ls())
```

Or, in RStudio, go to **Session > Restart R**.

# Import the 2016 & 2017 Divvy data

So far, we have only analyzed the trip data from the 4th quarter of 2017.

```
source("functions.R")
```

I wrote a function `read.divvy.data` to automate the reading and processing of all the downloaded Divvy data. It reads all the CSV files, then merges them into two data frames: one for the stations, and one for the trips.

# Import the 2016 & 2017 Divvy data

Choose which station and trip files to import:

```
stnfile    <- "Divvy_Stations_2017_Q3Q4.csv"  
tripfiles  <- Sys.glob("Divvy_Trips*.csv")
```

Variables `stnfile` and `tripfiles` contains the names of the files to be imported! they do not actually contain any data.

# Import the 2016 & 2017 Divvy data

This may take a minute to run, or longer if you have not installed the readr package.

```
divvy <- read.divvy.data(stnfile,tripfiles)
```

**Note:** If your computer does not have enough memory to load all the trip data, use only the Q1 trip files instead:

```
tripfiles <- Sys.glob("Divvy_Trips*Q1.csv")
```

What `read.divvy.data` does:

- Reads the Divvy station data from the CSV file.
- Reads the Divvy trip data from the CSV files.
- Combines the Divvy trip data into a single data frame.
- Takes additional steps to prepare the data.

# Import the 2016 & 2017 Divvy data

The output is a “list” containing two data frames. Let’s extract the data frames from the list:

```
names(divvy)
stations <- divvy$stations
trips    <- divvy$trips
rm(divvy)
head(stations)
head(trips)
nrow(trips)
```

- *Were more trips taken in 2016 or 2017?*
- *Which columns were converted to factors?*
- *What oddities do you notice from the summary?*

# Out first ggplot: a map of the Divvy stations

We will use the **ggplot2** package. It is a powerful (though not always intuitive) set of plotting functions that extend the base plotting functions in R.

```
install.packages ("ggplot2")
```

I also recommend the **cowplot** package, an extension to ggplot2 developed by Claus Wilke at UT Austin.

```
install.packages ("cowplot")
```

Load the ggplot2 and cowplot functions:

```
library(ggplot2)
```

```
library(cowplot)
```

# Plot station longitude vs. latitude

The “stations” data frame gives the geographic co-ordinates (latitude & longitude) for each station. With ggplot, we can create a station map from the “stations” data frame in only a few lines of code:

```
aes1 <- aes(x = longitude, y = latitude)
p      <- ggplot(stations, aes1)
print(p)
out <- geom_point()
p2  <- ggplot_add(out, p)
print(p2)
```

*What geographic features of Chicago are recognizable from this plot?*



# Adjusting the plot

Let's make a few adjustments to the plot:

```
out <- geom_point(shape = 21, fill = "darkblue",  
                  color = "white", size = 3)  
p3 <- ggplot_add(out, p)  
print(p3)
```

# Plotting contours instead of points

We can reuse our existing code, replacing the `geom_point` with a `geom_density_2d`, to create a very different plot:

```
out  <- geom_density_2d()  
p4   <- ggplot_add(out, p)  
print(p4)
```

# Use colors to highlight the largest stations

To do this, map the “dpcapacity” column to colour in the plot:

```
aes2 <- aes(x      = longitude,  
            y      = latitude,  
            color   = dpcapacity)  
p      <- ggplot(stations, aes2)  
out    <- geom_point()  
p <- ggplot_add(out, p)  
print(p)
```

The colour scale is not great, so let's improve it:

```
out <- scale_color_gradient2(low = "white",  
                             mid = "skyblue", high = "red", midpoint = 25)  
p <- ggplot_add(out, p)  
print(p)
```

*Where are the largest Divvy stations?*

# Scale stations by the number of departures

Next, let's add an additional piece of information to this visualization:

- Number of departures at each station, should (?) roughly correspond to population density.

To do this, we need to add a new column to the “stations” data frame containing the total number departures, which is calculated from the “trips” data frame:

```
counts <- table(trips$from_station_name)
```

Because we carefully prepared the data frame in `read.divvy.data`, station counts should be the same order as the stations. We can check this:

```
all(names(counts) == stations$name)
```

# Scale stations by the number of departures

Add these trip counts to the “stations” data frame:

```
stations$departures <- as.vector(counts)
head(stations)
```

Let's use this column in our new plot:

```
aes3 <- aes(x      = longitude,
            y      = latitude,
            size    = sqrt(departures))
p     <- ggplot(stations, aes3)
out   <- geom_point(shape = 21, fill = "blue",
                  color = "white")
p     <- ggplot_add(out, p)
print(p)
```

# How to save and share your plot

For exploratory analyses, GIF and PNG are great formats because the files are easy to attach to emails or webpages:

```
ggsave("station_map.png", p, dpi = 100)
```

For print or publication, save in a vector graphics format:

```
ggsave("station_map.pdf", p)
```

# Save your code & session state

This is a good time to save your session.

```
save.image("divvyanalysis.RData")
```

# Compare 2017 biking activity against 2016

Earlier, we observed an increase in trips from 2016 to 2017.  
Which stations experienced the largest increase?

- To examine this, we need to count trips separately for 2016 and 2017.
- Then we add these counts to the “stations” data frame.

We will use the `subset` and `table` to do this:

```
d1 <- subset(trips, start.year == 2016)
d2 <- subset(trips, start.year == 2017)
x1 <- table(d1$from_station_name)
x2 <- table(d2$from_station_name)
stations$dep.2016 <- as.vector(x1)
stations$dep.2017 <- as.vector(x2)
head(stations)
```



## Scatterplot of trips by station (2016 vs. 2017)

As before, now that we have prepared a data frame, plotting with ggplot is relatively straightforward:

```
aes3 <- aes(x = dep.2016, y = dep.2017)
p      <- ggplot(stations, aes3)
out    <- geom_point(shape = 20, size = 2)
p      <- ggplot_add(out, p)
print(p)
```

It is difficult to tell which stations had more trips in 2017—we need to compare against the  $x = y$  line.

```
out2 <- geom_abline(slope = 1, color = "skyblue",
                    linetype = "dashed")
p     <- ggplot_add(out2, p)
print(p)
```

*One station stands out because it has had a much larger increase in trips than other stations. What is this station?*

# Save your code & session state

Save your final results for safekeeping.

```
save.image("divvyanalysis.RData")
```

# ggplot: Take home points

- Creating sophisticated plots requires relatively little effort *provided the data are in the right form.*
- All plots in ggplot2 require these three elements:
  1. A data frame.
  2. An “aesthetic mapping” that declares how columns are mapped to plot features (axes, shapes, colors, *etc.*).
  3. A “geom”, short for “geometric object,” that specifies the type of plot.
- All plots are created by *adding layers*.

# Why data analysis in R?

- In R, a spreadsheet (“data frame”) is an object that can be inspected, manipulated and summarized with code.
- Therefore, we can write scripts to *automate* our data analyses.

# Parting thoughts

1. Always record your analysis steps in a file so you can reproduce them later.
2. Keep track of which packages (and the versions) you used with `sessionInfo()`.
3. Use packages—don't reinvent the wheel.
4. Email `help@rcc.uchicago.edu` for advice on using R on the RCC cluster.
5. Use “R Markdown” to document your analyses.
6. See the **workflowr** package for simplifying organizing & sharing of data analyses; e.g., **[stephenslab.github.io/wflow-divvy](https://stephenslab.github.io/wflow-divvy)**.
7. Thank you!